

Project Doc

First HBASE Image:

- First Download HBASE 2.5: <https://dlcdn.apache.org/hbase/2.5.11/hbase-2.5.11-bin.tar.gz>
- **Export Env Vars and Create HBASE Root Dir on HDFS**
 - `export HBASE_HOME=/data/hbase && export PATH=$HBASE_HOME/bin:$PATH`
 - `hdfs dfs -mkdir -p /hbase && hdfs dfs -chown hadoop:hadoop /hbase`
- copy hdfs-site.xml to the container
- doing this as multi-stage on hadoop image

Failover Tests :

- this is my active region servers

Region Servers

Base Stats						
Memory Requests Storefiles Compactions Replications						
ServerName	Start time	Last contact	Version	Requests Per Second	Num. Regions	
worker1.hbase_hadoop_cluster,16020,1747914326552	Thu May 22 11:45:26 GMT 2025	0 s	2.5.11	0	1	
worker2.hbase_hadoop_cluster,16020,1747917737696	Thu May 22 12:42:17 GMT 2025	0 s	2.5.11	0	1	
Total:2				0	2	

- now lets turn off one of them

```
PROBLEMS 49 OUTPUT DEBUG CONSOLE TERMINAL PORTS
Executing task: docker exec -it 02d5e29b4e4fda8e8aa7a68754ce796b33cd0235301d23cd21ab8538072d6e93 bash

hadoop@worker1:~$ jps
129 HRegionServer
11028 Jps
90 NodeManager
47 DataNode
hadoop@worker1:~$ kill 129
```

- now there is one region server , and has all regions bcs load balancer redirected the region to another region server

Region Servers

Base Stats	Memory	Requests	Storefiles	Compactions	Replications
ServerName	Start time	Last contact	Version	Requests Per Second	Num. Regions
worker2.hbase_hadoop_cluster,16020,1747917737696	Thu May 22 12:42:17 GMT 2025	1 s	2.5.11	0	3
Total:1				0	3

- now let's activate it using this command `hbase-daemon.sh start regionserver`

Region Servers

Base Stats	Memory	Requests	Storefiles	Compactions	Replications
ServerName	Start time	Last contact	Version	Requests Per Second	Num. Regions
worker1.hbase_hadoop_cluster,16020,1747928076759	Thu May 22 15:34:36 GMT 2025	1 s	2.5.11	0	1
worker2.hbase_hadoop_cluster,16020,1747917737696	Thu May 22 12:42:17 GMT 2025	2 s	2.5.11	0	2
Total:2				0	3

- now it's up and load balancer re-assigned the region to it

HMASTER Failover Test:

- i have 1 active and 2 backups

```
hbase:051:0> status
1 active master, 2 backup masters, 2 servers, 0 dead, 1.5000 average load
Took 0.3231 seconds
```

```
hbase:052:0> status 'detailed'
version 2.5.11
0 regionsInTransition
active master: hmaster1:16000 1747914317369
```

- HMaster1 is the Active Master
- now let's turn off it using kill command or turning off the container
- now my active is HMASTER 2 :

Backup Master hmaster2

Current Active Master: **hmaster2**

Tasks

Show All Monitored Tasks Show non-RPC Tasks Show All RPC Handler Tasks Show Active RPC Calls Show Client Operations

View as JSON

Start Time	Description	State	Status	Completion Time
Thu May 22 11:45:30 GMT 2025	Master startup	RUNNING (since 4hrs, 9mins, 48sec ago)	Initializing Master file system (since 1sec ago)	RUNNING

```
hbase:001:0> status
1 active master, 1 backup masters, 2 servers, 0 dead, 1.5000 average load
Took 1.1798 seconds
hbase:002:0>
```

- we can notice that i have 1 active and 1 backup and 1 dead
- now let's turn the container on:

Backup Masters

ServerName	Port	Start Time
hmaster1	16000	Thu May 22 15:59:42 GMT 2025
hmaster3	16000	Thu May 22 11:45:17 GMT 2025
Total:2		

-
- we can find it in the backup masters

```
hbase:002:0> status
1 active master, 2 backup masters, 2 servers, 0 dead, 1.5000 average load
Took 0.0344 seconds
```

Web Table:

HBase Table Schema: `webTable`

This HBase table is named `webTable` and is designed for efficient storage and retrieval of web content, metadata, and link structure. The schema uses salting, bloom filters, and finely tuned block settings to optimize performance for web-scale workloads.

Table Design Overview

Column Families

The table defines **four column families**, each with specific settings tailored for different access patterns:

1. `content`

- **Purpose:** Stores the main content of a webpage (e.g., HTML, text).
- **BLOOMFILTER:** `ROW` — optimized for lookups by row key.
- **BLOCKSIZE:** `65536` bytes (64 KB) — suited for larger values like full HTML pages.
- **BLOCKCACHE:** `true` — enables in-memory caching of blocks to improve read speed.
- **IN_MEMORY:** `true` — keeps data in memory for fast access (only viable if memory allows).
- **VERSIONS:** `1` — only the latest version of each cell is retained.

2. meta

- **Purpose:** Stores metadata such as titles, headers, timestamps.
 - **BLOOMFILTER:** ROW — fast row-level lookups.
 - **BLOCKSIZE:** 16384 bytes (16 KB) — smaller data blocks for lightweight metadata.
 - **BLOCKCACHE:** true — metadata is likely to be reused, so caching is beneficial.
 - **VERSIONS:** 1 — only the most recent metadata is needed.
-

3. outlinks

- **Purpose:** Stores links from the page to others (outbound links).
 - **BLOOMFILTER:** ROWCOL — optimized for access to specific link values by row and column.
 - **BLOCKSIZE:** 32768 bytes (32 KB) — moderate block size for lists of links.
 - **VERSIONS:** 1 — only the latest state of links is retained.
-

4. inlinks

- **Purpose:** Stores backlinks from other pages (who links to this page).
 - **BLOOMFILTER:** ROWCOL — enables precise lookup for specific incoming links.
 - **BLOCKSIZE:** 32768 bytes (32 KB).
 - **VERSIONS:** 1 .
-



Region Pre-splitting with Salts

```
['0!', '1!', '2!', '3!', '4!', '5!', '6!', '7!', '8!', '9!', 'a!', 'b!', 'c!',  
'd!', 'e!', 'f!']
```

```
create 'webTable',  
  {NAME => 'content',  
    BLOOMFILTER => 'ROW',  
    BLOCKSIZE => 65536,
```

```

BLOCKCACHE => true,
IN_MEMORY => true,
VERSIONS => 3,
TTL => 7776000},

{NAME => 'meta',
  BLOOMFILTER => 'ROW',
  BLOCKSIZE => 16384,
  BLOCKCACHE => true,
  VERSIONS => 1,
  TTL => 2147483647},

{NAME => 'outlinks',
  BLOOMFILTER => 'ROWCOL',
  BLOCKSIZE => 32768,
  VERSIONS => 2,
  TTL => 15552000},

{NAME => 'inlinks',
  BLOOMFILTER => 'ROWCOL',
  BLOCKSIZE => 32768,
  VERSIONS => 2,
  TTL => 15552000},

['0!', '1!', '2!', '3!', '4!', '5!', '6!', '7!', '8!', '9!', 'a!', 'b!',
'c!', 'd!', 'e!', 'f!']

```

- it's better to compress content with SNAPPY but need needs to be downloaded in HBASE lib

Why i used this table design:

- our table has heavy read and write but reads are more
- reading web content is much more that creating new content
- so i choose to use bloom filters , it might to delay write a bit but it will help us avoiding searching for data that doesn't exist
- we pre defined salted regions to avoid hotspotting
- we used this salting to help us designing our key

```
e!org.wikipedia.www/page-3
e!org.wikipedia.www/page-3
e!org.wikipedia.www/page-3
e!org.wikipedia.www/page-3

e!org.wikipedia.www/page-3

f!com.facebook.www/page-5

f!com.facebook.www/page-5
```

- this is keys sample , we reverse the web domain bcs all websites starts with www
- salting part , to make websites pages rows stored on the same region and same pages of the same domain stored with salt near to it's domain ,
 - for ex: google.com has salt 'e'
 - google.com/about has salt 'f'
 - but facebook for ex may get salt = 'a'

Queries:

Retrieve a page by exact row key

```
Took 0.0456 seconds
hbase:029:0> get 'webTable', 'e!org.wikipedia.www/page-3'
COLUMN                                CELL
content:html                          timestamp=2025-05-22T17:24:47.849, value=<html><h1>www.wikipedia.org /page-3</h1></html>
content:text                           timestamp=2025-05-22T17:24:47.849, value=welcome to www.wikipedia.org /page-3
meta:content_type                      timestamp=2025-05-22T17:24:47.849, value=text/html
meta:fetch_time                       timestamp=2025-05-22T17:24:47.849, value=1747934687839
meta:status                           timestamp=2025-05-22T17:24:47.849, value=200
outlinks:a!org.wikipedia.www/page-2   timestamp=2025-05-22T17:24:47.849, value=Previous Page
1 row(s)
Took 0.0132 seconds
hbase:030:0>
```

Scan all pages for a specific domain:

```
hbase:030:0> scan 'webTable', {FILTER => "PrefixFilter('e!org.wikipedia.www')"}
ROW COLUMN+CELL
e!org.wikipedia.www/page-3 column=content:html, timestamp=2025-05-22T17:24:47.849, value=<html><h1>www.wikipedia.org /page-3</h1></html>
e!org.wikipedia.www/page-3 column=content:text, timestamp=2025-05-22T17:24:47.849, value=welcome to www.wikipedia.org /page-3
e!org.wikipedia.www/page-3 column=meta:content_type, timestamp=2025-05-22T17:24:47.849, value=text/html
e!org.wikipedia.www/page-3 column=meta:fetch_time, timestamp=2025-05-22T17:24:47.849, value=1747934687839
e!org.wikipedia.www/page-3 column=meta:status, timestamp=2025-05-22T17:24:47.849, value=200
e!org.wikipedia.www/page-3 column=outlinks:a!org.wikipedia.www/page-2, timestamp=2025-05-22T17:24:47.849, value=Previous Page

1 row(s)
Took 0.0623 seconds
```

Find pages modified within a time range:

```
hbase:031:0> scan 'webTable', { FILTER => "SingleColumnValueFilter('meta', 'fetch_time', >=, 'binary:START_TIMESTAMP') AND " +
"SingleColumnValueFilter('meta', 'fetch_time', <=, 'binary:END_TIMESTAMP')"}
ROW COLUMN+CELL
c0!net.blog.tech/123 column=content:html, timestamp=2025-05-22T17:24:47.493, value=<html><p>Check out <a href="https://www.example.com">Example Inc</a></p></html>
c0!net.blog.tech/123 column=outlinks:com.example.www/, timestamp=2025-05-22T17:24:47.493, value=Example Inc

1 row(s)
Took 0.0789 seconds
```

Find inbound links to a specific page:

```
hbase:032:0> scan 'webTable', {FILTER => "QualifierFilter(=, 'substring:e!org.wikipedia.www/page-3')"}
ROW COLUMN+CELL
a!org.wikipedia.www/page-2 column=outlinks:e!org.wikipedia.www/page-3, timestamp=2025-05-22T17:24:47.842, value=Next Page

1 row(s)
Took 0.0456 seconds
```

Find pages containing specific text:

```
hbase:033:0> scan 'webTable', {FILTER => "ValueFilter(=, 'substring:Example')"}
ROW COLUMN+CELL
a0!com.example.www/ column=inlinks:net.blog.tech/123, timestamp=2025-05-22T17:24:47.500, value=Example Inc
c0!net.blog.tech/123 column=content:html, timestamp=2025-05-22T17:24:47.493, value=<html><p>Check out <a href="https://www.example.com">Example Inc</a></p></html>
c0!net.blog.tech/123 column=outlinks:com.example.www/, timestamp=2025-05-22T17:24:47.493, value=Example Inc

2 row(s)
Took 0.0249 seconds
```

Insert complete web page data (content, metadata, links):

```
put 'webTable', 'a0!com.example.www/', 'content:html', '<html>...</html>'
put 'webTable', 'a0!com.example.www/', 'meta:fetch_time', '1717020000000'
put 'webTable', 'a0!com.example.www/', 'meta:status', '200'
put 'webTable', 'a0!com.example.www/', 'outlinks:com.example.www/about',
'About Us'
```

```
put 'webTable', 'a0!com.example.www/', 'inlinks:net.blog.tech/123', 'Example Inc'
```

Retrieve a page by exact URL (row key):

```
hbase:039:0> get 'webTable', 'a0!com.example.www/'
COLUMN                                CELL
content:html                          timestamp=2025-05-22T17:34:18.587, value=<html>...</html>
content:text                          timestamp=2025-05-22T17:24:47.477, value=Welcome to our site. About Us.
inlinks:net.blog.tech/123             timestamp=2025-05-22T17:34:39.154, value=Example Inc
meta:content_type                    timestamp=2025-05-22T17:24:47.477, value=text/html
meta:fetch_time                      timestamp=2025-05-22T17:34:23.326, value=1717020000000
meta:status                          timestamp=2025-05-22T17:34:31.395, value=200
outlinks:com.example.www/about       timestamp=2025-05-22T17:34:35.923, value=About Us
t
1 row(s)
Took 0.0185 seconds
```

Update a page's content and metadata:

```
put 'webTable', 'a0!com.example.www/', 'content:html', '<html><h1>Updated Content</h1></html>'
put 'webTable', 'a0!com.example.www/', 'meta:fetch_time', '1717030000000'
```

Delete a page and all its information:

```
hbase:042:0> deleteall 'webTable', 'a0!com.example.www/'
Took 0.1167 seconds
hbase:043:0> 
```

Find pages with titles containing specific keywords:

```
hbase:043:0> scan 'webTable', {FILTER => "ValueFilter(=, 'substring:Keyword')"}
ROW                                COLUMN+CELL
0 row(s)
Took 0.0129 seconds
```

List pages with specific HTTP status codes:


```

hbase:044:0> scan 'webTable', {FILTER => "SingleColumnValueFilter('meta', 'status', =, 'binary:404')"}
ROW          COLUMN+CELL
b0!com.example.www/about  column=content:html, timestamp=2025-05-22T17:24:47.485, value=<html><h2>About</h2><p>
Our company info</p></html>
b0!com.example.www/about  column=inlinks:com.example.www/, timestamp=2025-05-22T17:24:47.485, value=Home
b0!com.example.www/about  column=meta:fetch_time, timestamp=2025-05-22T17:24:47.485, value=1717020001000
c0!net.blog.tech/123      column=content:html, timestamp=2025-05-22T17:24:47.493, value=<html><p>Check out <a h
ref="https://www.example.com">Example Inc</a></p></html>
c0!net.blog.tech/123      column=outlinks:com.example.www/, timestamp=2025-05-22T17:24:47.493, value=Example In
c
2 row(s)
Took 0.0181 seconds

```

```

hbase:045:0> scan 'webTable', {FILTER => "SingleColumnValueFilter('meta', 'status', >=, 'binary:400')"}
ROW          COLUMN+CELL
b0!com.example.www/about  column=content:html, timestamp=2025-05-22T17:24:47.485, value=<html><h2>About</h2><p>
Our company info</p></html>
b0!com.example.www/about  column=inlinks:com.example.www/, timestamp=2025-05-22T17:24:47.485, value=Home
b0!com.example.www/about  column=meta:fetch_time, timestamp=2025-05-22T17:24:47.485, value=1717020001000
c0!net.blog.tech/123      column=content:html, timestamp=2025-05-22T17:24:47.493, value=<html><p>Check out <a h
ref="https://www.example.com">Example Inc</a></p></html>
c0!net.blog.tech/123      column=outlinks:com.example.www/, timestamp=2025-05-22T17:24:47.493, value=Example In
c
2 row(s)
Took 0.0103 seconds

```

Find pages modified after a specific date:

```

hbase:046:0> scan 'webTable', {FILTER => "SingleColumnValueFilter('meta', 'fetch_time', >, 'binary:TIMESTAMP')"}
ROW          COLUMN+CELL
c0!net.blog.tech/123      column=content:html, timestamp=2025-05-22T17:24:47.493, value=<html><p>Check out <a h
ref="https://www.example.com">Example Inc</a></p></html>
c0!net.blog.tech/123      column=outlinks:com.example.www/, timestamp=2025-05-22T17:24:47.493, value=Example In
c
1 row(s)
Took 0.0218 seconds

```

Manual purge for outdated content:

```

Took 0.0255 seconds
hbase:050:0> delete 'webTable', 'a0!com.example.www/', 'content:html', 1680000000000 # delete specific timestamp ve
rsion
Took 0.0148 seconds
hbase:051:0> deleteall 'webTable', 'a0!com.example.www/' # delete entire row
Took 0.0073 seconds

```

Retrieve latest N versions of content:

```

hbase:065:0> get 'webTable', 'a0!com.example.www/', {COLUMN => 'content:html', VERSIONS => 3}
COLUMN      CELL
0 row(s)
Took 0.0042 seconds

```